YICHI ZHANG

+1(734) 800-6032 ◇ yz1636@dimacs.rutgers.edu ◇ Website: https://yichiz97.github.io/
CoRE 417, Computer Science Department, Rutgers University
96 Frelinghuysen Rd, Piscataway, NJ 08854, USA

EDUCATION

University of Michigan, Ann Arbor (Umich) Sep 2019 - Dec 2024

- Ph.D. in Information
- Advisor: Grant Schoenebeck
- Thesis: Incentivizing Effort and Honesty for High-quality Information

Shanghai Jiao Tong University (SJTU) Aug 2015 - Jun 2019

- B.S. in Electronic Science and Engineering
- Advisors: Xinbing Wang and Luoyi Fu

PROFESSIONAL EXPERIENCE

Postdoctoral Associate, *The Center for Discrete Mathematics and Theoretical Computer Science (DIMACS), Rutgers University* Sep 2024 - Present

- Hosts: David Pennock and Lirong Xia

Graduate Student Research Assistant, *School of Information, University of Michigan* Sep 2019 - Aug 2024

Research Intern, Department of Computer Science, UCLA Jul 2018 - Sep 2018

- Mentor: Mario Gerla

Algorithm Engineer Intern, YITUTech Feb 2019 - May 2019

RESEARCH INTERESTS

Design theory-grounded, verification-independent, manipulation-resistant evaluation metrics that elicit high-effort human feedback, quantify data quality, and steer AI development.

- Methodology: combine information elicitation, game theory, and mechanism design with machine learning/LLMs.
- Applications: crowdsourcing, peer grading/review, and LLM alignment & ensemble.

WORKING PAPERS

Good Enough? Evaluating Peer and AI Grading via A TA Benchmark
Sanzeed Anwar*,**Yichi Zhang***, Noah Burrell, and Grant Schoenebeck
Under Review at AAAI 2026

Mixture of Complementary Agents for Robust LLM Ensemble
Yichi Zhang, Kevin Lu, Yuang Zhang, Jie Gao, Lirong Xia, and Fang-Yi Yu
Under Review at ICLR 2026

Supervised Fine-Tuning on Ambiguous Preference Pairs Boosts LLM Alignment
Jinlong Pang, Zhaowei Zhu, Na Di, **Yichi Zhang**, Yaxuan Wang, Chen Qian, Yang Liu
Under Review at ICLR 2026

From Crowds to Codes: Can Adaptive Peer Review Help?
Xingbo Wang, Fang-Yi Yu, **Yichi Zhang** (alphabetically ordered)
Under Review at ICLR 2026

JOURNAL SUBMISSIONS

A System-Level Analysis of Conference Peer Review
Yichi Zhang, Fang-Yi Yu, Grant Schoenebeck, and David Kempe
Major Revision at Operations Research

CONFERENCE PUBLICATIONS

Evaluating LLM-Corrupted Crowdsourcing Data Without Verifications
Yichi Zhang*, Jinlong Pang*, Zhaowei Zhu, and Yang Liu
In Proceedings of the 39th Annual Conference on Neural Information Processing Systems (NeurIPS 2025)
[https://arxiv.org/abs/2506.06991]

Stochastically Dominant Peer Prediction
Yichi Zhang, Shengwei Xu, David Pennock, and Grant Schoenebeck
In Proceedings of the 39th Annual Conference on Neural Information Processing Systems (NeurIPS 2025)
[https://arxiv.org/abs/2506.02259]

Eliciting Informative Text Evaluations with Large Language Models
Yuxuan Lu, Shengwei Xu, **Yichi Zhang**, Yuqing Kong, and Grant Schoenebeck
In Proceedings of the 25th ACM Conference on Economics and Computation (EC 2024)
[https://arxiv.org/abs/2405.15077]

Spot Check Equivalence: an Interpretable Metric for Information Elicitation Mechanisms
Shengwei Xu, **Yichi Zhang**, Paul Resnick, and Grant Schoenebeck
In Proceedings of the 33nd Annual World Wide Web Conference (WWW 2024)
[https://arxiv.org/abs/2402.13567]
(Oral presentation)

Eliciting Honest Information From Authors Using Sequential Review
Yichi Zhang, Grant Schoenebeck, and Weijie Su
In Proceedings of the 38th Annual AAAI Conference on Artificial Intelligence (AAAI 2024)
[https://arxiv.org/abs/2311.14619]

Multi-task Peer Prediction Under Task-Dependent Strategies
Yichi Zhang and Grant Schoenebeck
In Proceedings of the 32nd Annual World Wide Web Conference (WWW 2023)
[https://dl.acm.org/doi/abs/10.1145/3543507.3583292]

High-Effort Crowds: Limited Liability Via Tournaments
Yichi Zhang and Grant Schoenebeck
In Proceedings of the 32nd Annual World Wide Web Conference (WWW 2023)
[https://dl.acm.org/doi/abs/10.1145/3543507.3583334]

A System-Level Analysis of Conference Peer Review
Yichi Zhang, Fang-Yi Yu, Grant Schoenebeck, and David Kempe

In Proceedings of the 23rd ACM Conference on Economics and Computation (EC 2022)
[https://arxiv.org/abs/2303.09020]

Information Elicitation From Rowdy Crowds
Grant Schoenebeck, Fang-Yi Yu, and **Yichi Zhang** (alphabetically ordered) *In Proceedings of the 30th Annual World Wide Web Conference* (WWW 2021)
[https://dl.acm.org/doi/abs/10.1145/3442381.3449840]

POSITION PAPERS

Strategic Foundation Models
Denizalp Goktas, Amy Greenwald, Takayuki Osogami, Roma Patel, Kevin Leyton-Brown, Grant Schoenebeck, Daphne Cornelisse, Constantinos Daskalakis, Ian Gemp, John Horton, David C Parkes, David M Pennock, Arjun Prakash, Sai Srivatsa Ravindranath, Max Olan Smith, Gokul Swamy, Eugene Vinitsky, Segev Wasserkrug, Michael Wellman, Jibang Wu, Haifeng Xu, Jiayao Zhang, **Yichi Zhang**, Sadie Zhao, Quanyan Zhu
[https://hal.science/hal-04925309]

INVITED TALKS

Evaluating LLM-Corrupted Crowdsourcing Data Without Verification

- Stanford University, the EC Workshop on Human–Algorithm Collaboration July 2025

Peer Prediction on the Move: From Expected Score to Score Distribution

- Princeton University, Mechanism Design Group November 2024

High-Effort Crowds: Limited Liability Via Tournaments

- Rutgers University, the DIMACS Workshop on Forecasting October 2024

Eliciting Honest Information From Authors Using Sequential Review

- Princeton University, Mechanism Design Group October 2024
- University of Massachusetts, Amherst, Computer Science Theory Seminar October 2024
- Harvard University, EconCS seminar January 2024

Improving Conference Review Via Mechanism Design

- Renmin University of China, Gaoling School of Artificial Intelligence August 2023
- Peking University, CFCS seminar July 2023
- University of Michigan, DSCSS seminar April 2023
- University of Pennsylvania, Wharton Statistics and Data Science March 2023
- Drexel University, Computer Science Department March 2023

TEACHING

Guest lecturer, Rutgers **Fall 2025**
Courses: Trustworthy Statistical Learning (Large Language Model) (STATS 656)
Instructor: Linjun Zhang

Guest lecturer, Rutgers **Fall 2025**
Courses: Introduction to AI (CS 440)
Instructor: Lirong Xia

Teaching assistant (GSI), Umich **Winter 2022**
Courses: Big Data Analysis (SI 699)
Instructor: Misha Teplitskiy

Teaching assistant (GSI), Umich **Fall 2021**
Course: Deep Learning (SIADS 642)
Instructor: Paramveer Dhillon
Course: Network Analysis (SIADS 642)
Instructor: Daniel Romero

AWARDS

- ICSSI Travel Award. 2024
- Rackham Conference Travel Grant, University of Michigan. 2023, 2024
- The Web Conference Student Scholarship. 2021
- Nominee for the Rackham International Student Fellowship, UMSI. 2021
- Outstanding Graduate of Shanghai Jiao Tong University. 2019
- EIC Education Scholarship (top 5%). 2018
- Samsung Scholarship (top 3%). 2017

SERVICE

Journal Reviewer:

- Transactions on Machine Learning Research **(TMLR)**: 2025

Conference Reviewer/PC member:

- International Conference on Learning Representations **(ICLR)**: 2025–2026
- The Annual Conference on Neural Information Processing Systems **(NeurIPS)**: 2024–2025
- The Web Conference **(WWW)**: 2022–2025
- The ACM Conference on Economics and Computation **(EC)**: 2025
- Conference on Web and Internet Economics **(WINE)**: 2023–2024
- International Conference on Computational Social Science **(IC2S2)**: 2024

Organizer:

- The Annual Workshop on Incentives in Academia, jointly organized with EC 2024, 2025.